

Chapter 1

Algorithmic Trading and Market Microstructure

1.1 Introduction

Every strategy considered so far in this book, from the DuPont-decomposed valuation of Chapter 4 to the mean-variance portfolios of Chapter 12, has treated buying and selling an asset as instantaneous and free, an assumption made deliberately in each of those chapters so that the underlying financial or statistical idea could be presented without an additional layer of execution detail. Chapter 3 already showed this is not quite right: every trade costs at least half the bid-ask spread, and a large enough order walks through several levels of the order book at a progressively worse price. This chapter treats execution itself as a subject worth studying carefully, not a frictionless afterthought: how a trading signal is generated, how an order is actually worked in the market, how a strategy's historical performance should be evaluated, and how the market microstructure concepts of Chapter 3 shape every one of those decisions.

Several worked examples run through the chapter, each grounded in material introduced earlier in this book. A moving-average crossover strategy, backtested on a short synthetic price series, illustrates how a trading signal becomes a sequence of positions and, ultimately, a profit-and-loss series. The limit-order-book walk from Chapter 3 is extended into a comparison between aggressive, single-shot execution and a patient, multi-slice alternative, motivating the standard TWAP and VWAP execution benchmarks, and into a regression-based estimate of Kyle's lambda that gives price impact a formal economic interpretation.

The market-making intuition of Chapter 3 is developed into a fully worked Avellaneda-Stoikov inventory example, showing exactly how and why a market maker's quotes should skew as its position grows.

The pairs-trading setup introduced conceptually in Chapter 7 is extended into an actual backtested trade, entered and exited using the same z -score rule that chapter only described qualitatively, and then generalized into a multi-stock statistical arbitrage example built on Chapter 12's factor-regression machinery.

As throughout this book, every numerical claim below was computed in Python first and is reproduced exactly in the companion notebook.

This chapter sits at a specific point in the book's overall structure that is worth making explicit. Chapters 6 and 7 supplied the statistical and machine learning tools capable of generating a trading signal in the first place; Chapter 12 supplied the portfolio-construction logic for combining many assets and signals into a coherent overall position; and this chapter supplies the missing link between the two, everything that happens after a model or a portfolio optimizer has decided what it wants to hold, and before that decision becomes a completed, real-world trade at a real, and imperfect, price.

Skipping this link, treating execution as instantaneous and free, is precisely the simplifying assumption every earlier chapter in this book has made for clarity, and this chapter exists specifically to show what that assumption costs when it is finally relaxed.

1.2 A Taxonomy of Trading Strategies and Signals

Suppose two traders both notice the exact same pattern in a stock's recent price history but reach opposite conclusions about what to do with it: one buys, betting a recent rise will continue; the other sells short, betting the same rise has already gone too far and must reverse. Both can be running a perfectly coherent strategy; they simply disagree about the economic source of the pattern each believes is in front of them.

Algorithmic trading strategies are usually organized by exactly this distinction, the economic source of the pattern they exploit, and it is worth naming the major categories before working through any one of them in detail. Momentum strategies bet that a recent price trend will continue, on the theory that information diffuses gradually through a market rather than being reflected in prices instantly, so that a stock which has recently risen is more likely than not to keep rising over some short additional horizon. Mean-reversion strategies bet the opposite: that a price (or a relationship between two prices, as in the pairs-trading strategy developed later in this chapter) has moved too far from some fundamental or statistical anchor and will revert toward it, an application of exactly the stationarity concepts Chapter 7 developed at length. Market-making strategies, discussed further in Section 1.9, do not take a directional view on price movement at all; they profit from repeatedly capturing the bid-ask spread while managing the inventory risk that results from doing so. Statistical arbitrage strategies generalize mean reversion and pairs trading to many assets simultaneously, often using the factor models of Chapter 12 to identify a stable, model-implied relationship worth betting will hold.

Every one of these strategy types ultimately produces the same basic output: a signal, a number (or simply a direction) indicating whether and how much to buy or sell at a given point in time. Converting that signal into an actual position, and that position into actual executed trades, is the specific concern of most of this chapter, and it is a step with its own economics, costs, and risks that are entirely separate from whether the underlying signal is any good in the first place.

1.3 A Worked Example: A Moving-Average Crossover Strategy

A simple moving-average crossover strategy, one of the oldest and most transparent momentum strategies, holds a long position whenever a short-window moving average is above a longer-window moving average, and is flat otherwise, on the theory that the short average crossing above the long average signals a strengthening upward trend. Table 1.1 applies a 3-day short average and a 6-day long average to twelve days of hypothetical prices.

Day	Price	MA(3)	MA(6)	Signal	Strategy Return
1	100	—	—	0	—
2	101	—	—	0	0.00%
3	103	101.33	—	0	0.00%
4	106	103.33	—	0	0.00%
5	110	106.33	—	0	0.00%
6	115	110.33	105.83	1	0.00%
7	119	114.67	109.00	1	3.48%
8	118	117.33	111.83	1	−0.84%
9	114	117.00	113.67	1	−3.39%
10	109	113.67	114.17	0	−4.39%
11	104	109.00	113.17	0	0.00%
12	100	104.33	110.67	0	0.00%

Table 1.1: A 3-day/6-day moving-average crossover strategy

The signal turns on (long) at day 6, when the 3-day average first exceeds the 6-day average, and turns off (flat) at day 10, when it first falls back below. Crucially, the strategy return in each row uses the *previous* day's signal applied to the *current* day's price return, since a signal computed from today's closing price cannot actually be traded until at least the next available price, exactly the point-in-time discipline Chapter

6 emphasized when warning against data leakage; using today's signal against today's own return would silently look ahead.

Compounding the strategy returns in the table gives a cumulative return of -5.22% over the twelve days, while simply buying and holding the asset from day 1 to day 12 (which starts and ends at a price of 100) returns exactly 0%. In this particular, deliberately illustrative sample, the crossover strategy actually underperforms a naive buy-and-hold position: it captures part of the days-6-through-8 uptrend but stays long one day too long into the days-9-and-10 reversal, precisely the whipsaw risk that afflicts every momentum strategy during a trend reversal. This is an intentionally honest example, not a cherry-picked success story: a real backtest evaluation, developed further in Section 1.13, must be prepared to find that a plausible-sounding strategy simply does not work on the available data.

1.4 Order Types and the Mechanics of Execution

Chapter 3 introduced the market order (executed immediately at the best available price, walking through the book if necessary) and the limit order (executed only at a specified price or better, potentially never filling at all). Converting a trading signal like the one in Section 1.3 into an actual position requires choosing among these order types, and the choice involves a direct tradeoff: a market order guarantees execution but accepts whatever price impact results, while a limit order controls the execution price but risks not being filled at all if the market moves away before the order reaches the front of the queue. A stop order, a third common type, becomes a market order automatically once a specified trigger price is reached, commonly used to implement a predetermined exit rule (a stop-loss) without requiring continuous manual monitoring of the position.

Beyond the order type itself, an algorithmic execution system must decide how to route an order, since modern equity markets consist of many competing exchanges and alternative trading venues rather than a single centralized order book, a fragmentation Chapter 3 touched on in its discussion of Regulation NMS. A smart order router evaluates the available liquidity and price across every venue simultaneously and splits or directs an order to capture the best aggregate execution, a considerably more complex problem than the single-order-book walk introduced in Chapter 3, but built on exactly the same underlying economics of price and available size at each level.

A large order also raises a genuine tension that the order types above do not fully resolve: posting the full order as a visible limit order reveals the trader's intentions to every other market participant, who can then adjust their own behavior in response, for example by front-running the anticipated remaining demand. Iceberg orders address this by displaying only a small portion of the total order size at any given time, automatically revealing more once the visible portion is filled, so that the order book shows only a fraction of the true remaining size. Dark pools take this further still, entirely private trading venues that match buy and sell orders without displaying any pre-trade price or size information to the broader market, only reporting completed trades after the fact. Both mechanisms trade away pre-trade transparency, one of the market-quality benefits Chapter 3 attributed to well-functioning exchanges, in exchange for reduced information leakage and, in principle, lower market impact for large orders; this is a genuine tradeoff rather than a free improvement, since a market with a large fraction of its volume trading in dark, non-displayed venues can make the lit, publicly visible order book a less complete and less reliable picture of overall supply and demand, a concern regulators have studied closely as dark-pool trading volume has grown.

1.5 Market Impact and the Cost of Trading Fast

Chapter 3's limit-order-book example showed that a 250-share market order, larger than the 100 shares available at the best ask of \$50.10, walked through a second price level and executed at an average price of \$50.13, a price impact of about 6 basis points relative to the best ask. This cost arises entirely from the size of the order relative to the liquidity available at the time; it is not a fee charged by anyone, but the direct consequence of consuming the available supply at each successive price level.

An alternative to executing the entire order at once is to break it into smaller pieces and space them out over time, allowing the order book to replenish between slices. Suppose the same 250-share order is instead split into five 50-share slices, each small enough to execute entirely within the 100 shares available at the




Figure 1.1: The path from a trading signal to a completed, evaluated trade: each stage introduces its own decisions and costs, and transaction cost analysis at the end feeds back into how future signals and execution choices are made.

best ask in Chapter 3’s example, and suppose the best ask remains at \$50.10 between slices (a simplifying assumption; in reality, the book may or may not fully replenish, and the price may drift for reasons entirely unrelated to this trader’s own activity).

This patient approach achieves an average execution price of exactly \$50.10, eliminating the price impact entirely, but at a cost the single-shot execution does not bear: over the time it takes to work five separate slices, the underlying price can move for entirely unrelated reasons, a risk called timing risk or execution risk. The aggressive, single-shot strategy pays a known, certain price-impact cost immediately; the patient, multi-slice strategy pays an expected impact cost of zero but accepts an uncertain, and potentially much larger or smaller, timing-risk cost. Neither strategy is universally superior; which one a trader should prefer depends on the size of the order relative to available liquidity, the asset’s volatility, and how urgently the position needs to be established, exactly the tradeoff formalized in Section 1.8.

1.6 The Kyle Model: A Formal Model of Price Impact

Chapter 3’s order-book walk demonstrates price impact mechanically, one order consuming successive price levels, but does not explain *why* a market maker sets prices that respond to order flow in the first place. Kyle’s (1985) classic model supplies the economic explanation: a market maker cannot distinguish an order coming from an informed trader (who trades because they know something about the asset’s true value) from one coming from a liquidity trader (who trades for reasons unrelated to information, such as needing cash), and must set a single price schedule that, on average, breaks even against this mixture. The market maker’s rational response is to move the price in proportion to the net signed order flow it observes,

$$\Delta P = \lambda \times (\text{Net Order Flow}) + \varepsilon,$$

where net order flow is buy volume minus sell volume over some interval and λ (Kyle’s lambda) measures the price impact per unit of net order flow, exactly the same linear-regression structure used throughout this book, here estimated by regressing observed price changes on observed net order flow.

Table 1.2 shows six intervals of net order flow (in thousands of shares) and the resulting price change (in cents).

Interval	1	2	3	4	5	6
Net order flow (000s)	5	−3	8	−6	2	4
Price change (cents)	10	−7	15	−11	3	9

Table 1.2: Net order flow and price change for estimating Kyle’s lambda

Applying the same OLS formula used since Chapter 6,

$$\hat{\lambda} = \frac{\sum_t (\text{Flow}_t - \overline{\text{Flow}})(\Delta P_t - \overline{\Delta P})}{\sum_t (\text{Flow}_t - \overline{\text{Flow}})^2} \approx 1.95 \text{ cents per thousand shares,} \quad R^2 \approx 0.99.$$

A larger estimated λ indicates a market in which prices move more sharply for a given amount of net buying or selling pressure, which Kyle's model interprets as a market in which informed trading is a larger fraction of total order flow, or in which liquidity (the depth available to absorb order flow without moving the price, exactly the order-book depth introduced in Chapter 3) is thinner. This gives price impact, treated so far in this chapter as a mechanical consequence of a specific order book, a deeper economic interpretation: it is compensation the market maker requires for the risk of unknowingly trading against better-informed counterparties, the same adverse-selection component of the spread Chapter 3 first introduced, now expressed as a single estimable slope coefficient rather than only a qualitative concept.

1.7 TWAP and VWAP: Basic Execution Algorithms

The patient, multi-slice approach in Section 1.5 is a simple version of the time-weighted average price (TWAP) algorithm: split an order into equal-sized slices spread evenly across a time horizon, without reference to how trading volume is actually distributed over that horizon. Table 1.3 shows five hypothetical intervals with market prices and total trading volume.

Interval	1	2	3	4	5
Price	\$50.00	\$50.05	\$50.10	\$50.08	\$50.12
Market volume	10,000	12,000	8,000	15,000	9,000

Table 1.3: Prices and market volume over five execution intervals

A trader executing a 1,000-share order via TWAP trades 200 shares in each interval regardless of volume, achieving an average execution price equal to the simple average of the five prices,

$$\text{TWAP execution price} = \frac{50.00 + 50.05 + 50.10 + 50.08 + 50.12}{5} = \$50.07.$$

The volume-weighted average price (VWAP) benchmark, by contrast, weights each interval's price by that interval's share of total market volume,

$$\text{VWAP} = \frac{\sum_t P_t V_t}{\sum_t V_t} = \frac{50.00(10,000) + \dots + 50.12(9,000)}{10,000 + 12,000 + 8,000 + 15,000 + 9,000} \approx \$50.068,$$

close to, but not identical to, the naive TWAP execution price, a difference of about 0.4 basis points in this example. The gap arises because interval 4, with the largest volume (15,000 shares) in the table, happened to occur at a below-average price (\$50.08); a VWAP-following execution algorithm would have traded proportionally more in that favorable interval than the equally-weighted TWAP schedule did, capturing a slightly better average price. In practice, a VWAP execution algorithm typically trades according to a historical intraday volume profile, since the actual volume for the current day is not known in advance. It attempts to match the market's own volume distribution rather than the simpler, uniform TWAP schedule, on the theory that trading in proportion to when the market is naturally most liquid minimizes the executing trader's own footprint and price impact.

1.8 Optimal Execution: Balancing Impact and Timing Risk

Section 1.5's comparison between an aggressive, single-shot execution and a patient, multi-slice alternative is a simplified version of the optimal execution problem formalized by Almgren and Chriss: given an order of a fixed size that must be completed by a fixed deadline, how should it be split over time to minimize the combination of expected market impact cost and the variance introduced by timing risk? Market impact cost is typically modeled as increasing in the trading rate, since trading faster within any given interval consumes more of the available liquidity at each price level, exactly as in Chapter 3's order-book walk. Timing risk, by contrast, is modeled as increasing in the total time the order remains unexecuted, since a longer execution horizon exposes the remaining unexecuted quantity to more periods of potentially adverse price movement,

using the same variance-grows-with-horizon logic Chapter 7 developed for multi-step forecasting. The key output of this framework is not a single number but an efficient frontier directly analogous to the risk-return efficient frontier of Chapter 12. For a given level of acceptable timing risk, there is a specific optimal trading trajectory that minimizes expected execution cost, typically front-loaded, trading more aggressively early and tapering off, rather than either fully aggressive or perfectly uniform. A more risk-averse trader, one who weights the variance of the outcome more heavily relative to its expected cost, will choose a faster, more front-loaded execution than a less risk-averse trader facing the identical order and market conditions. This formalizes, with actual mathematical structure, the intuitive tradeoff Section 1.5 described only qualitatively: neither the purely aggressive nor the purely patient extreme is generally optimal, and the right answer depends on quantities, market impact sensitivity and price volatility, that must themselves be estimated from data, with all the estimation-error caveats Chapter 12 raised about any input estimated from a limited historical sample.

1.8.1 A Worked Example: The Almgren-Chriss Optimal Trajectory

Almgren and Chriss formalize the tradeoff above by minimizing a risk-adjusted total cost, expected temporary-impact cost plus λ times its variance from timing risk, where $\lambda \geq 0$ is the trader's risk-aversion parameter,

$$\min_{x_1, \dots, x_{N-1}} \underbrace{\eta \sum_{j=1}^N \frac{(x_{j-1} - x_j)^2}{\tau}}_{\text{expected impact cost}} + \lambda \underbrace{\sigma^2 \tau \sum_{j=1}^N x_j^2}_{\text{timing-risk variance}},$$

where x_j is the number of shares still unexecuted after interval j (so $x_0 = X$, the full order, and $x_N = 0$, fully executed), τ is the length of each interval, η is the temporary-impact cost coefficient, and σ^2 is the per-interval return variance. The closed-form solution to this minimization is

$$x_j = X \frac{\sinh(\kappa(T - t_j))}{\sinh(\kappa T)}, \quad \kappa = \sqrt{\frac{\lambda \sigma^2}{\eta}},$$

with $T = N\tau$ the total execution horizon; κ is the single parameter that governs how front-loaded the optimal trajectory is, growing with risk aversion λ and volatility σ and shrinking with the impact cost η . The trajectory is stated here without proof; this book's companion technical note on the Almgren-Chriss model derives it from the discrete first-order condition, which turns out to be a linear second-order difference equation, and verifies the two limiting cases that make the result trustworthy: uniform, TWAP-style trading as risk aversion vanishes, and immediate liquidation as it grows without bound.

Applying this to Section 1.5's 250-share order split across the same $N = 5$ intervals, with illustrative parameters $\sigma = 2\%$ per interval, $\eta = 0.01$, and a modest risk aversion of $\lambda = 0.5$ (so $\kappa \approx 0.1414$), the optimal trajectory leaves

$$\begin{aligned} \text{Unexecuted shares: } x &= (250, 194.2, 142.4, 93.4, 46.2, 0), \\ \text{Traded per interval: } &(55.8, 51.9, 49.0, 47.1, 46.2), \end{aligned}$$

mildly front-loaded relative to the perfectly uniform 50-share-per-interval TWAP schedule of Section 1.5. Raising risk aversion tenfold to $\lambda = 5$ (so $\kappa \approx 0.4472$) sharpens this front-loading considerably,

$$\text{Traded per interval: } (92.8, 60.9, 41.3, 30.1, 25.0),$$

nearly double the first-interval quantity of the low-risk-aversion trajectory: a more risk-averse trader accepts a higher expected impact cost from trading more aggressively up front specifically to shrink the variance contributed by however much of the order remains exposed to further price moves. Figure 1.2 plots both optimal trajectories against the uniform TWAP baseline, making the front-loading pattern, and how much more pronounced it becomes as λ rises, directly visible.

Scoring all three strategies, this front-loaded optimal trajectory, the uniform TWAP schedule, and the fully aggressive single-shot execution of Section 1.5, on the identical risk-adjusted objective at $\lambda = 0.5$ makes




Figure 1.2: Almgren-Chriss optimal unexecuted-share trajectories at two risk-aversion levels, against the uniform TWAP baseline.

the tradeoff numerically concrete:

$$\begin{aligned}
 \text{Optimal trajectory:} & \text{ cost} = 139.4, \\
 \text{Uniform TWAP:} & \text{ cost} = 140.0, \\
 \text{Fully aggressive single-shot:} & \text{ cost} = 625.0,
 \end{aligned}$$

the TWAP schedule barely worse since $\lambda = 0.5$ is only mildly risk-averse here, and the fully aggressive execution over four times worse, since it eliminates timing-risk variance entirely but pays a quadratically larger impact cost for trading the entire order in one interval. This is exactly the efficient-frontier logic described qualitatively above, made concrete: the fully aggressive and perfectly uniform schedules are both feasible points, but neither is optimal once risk aversion is weighed formally against expected cost, and the specific front-loaded shape of the true optimum, more pronounced the more risk-averse the trader, is not a generic property of “spreading an order out” but the precise output of this closed-form solution.

1.9 Inventory Models and Market Making

Chapter 3 introduced market making conceptually: a market maker profits from the bid-ask spread by standing ready to both buy and sell, but bears inventory risk from holding a non-zero position while waiting for an offsetting order to arrive. The Avellaneda-Stoikov model formalizes this by having the market maker continuously adjust both quotes based on current inventory, not just on the underlying asset’s fair value. A market maker who has accumulated a large long position, from selling less than buying, should skew both the bid and ask downward: reducing the bid makes it relatively less attractive for other traders to sell to the market maker, and lowering the ask makes it relatively more attractive to buy from it. That actively works the existing inventory back toward zero rather than passively waiting for random order flow to do so.

The model formalizes this skew through a reservation price, the price at which the market maker would be indifferent to a marginal trade given its current inventory, which sits below the true mid price when the market maker is long and above it when short,

$$r = s - q \gamma \sigma^2 (T - t),$$

where s is the current mid price, q is the market maker’s current inventory (positive if long, negative if short), γ is a risk-aversion parameter, σ^2 is the asset’s return variance, and $T - t$ is the time remaining in

the market maker's trading horizon. The optimal bid and ask are then centered on this reservation price, not on the mid price itself, and separated by an optimal spread

$$\delta = \gamma \sigma^2 (T - t) + \frac{2}{\gamma} \ln \left(1 + \frac{\gamma}{\kappa} \right),$$

where κ governs how sensitively order arrival rates respond to how competitively the market maker is quoting. Consider a market maker with mid price $s = \$100$, currently long $q = 5$ units, $\gamma = 0.1$, $\sigma = 2$, one full unit of time remaining ($T - t = 1$), and $\kappa = 1.5$. The reservation price is

$$r = 100 - 5(0.1)(2^2)(1) = \$98.00,$$

two dollars below the mid price, reflecting the market maker's desire to reduce its long position. Both halves of the spread formula above are quoted rather than derived; this book's companion technical note on the Avellaneda-Stoikov model obtains each exactly and separately, the reservation price from exponential (CARA) utility applied to a normally distributed terminal mark-to-market value, and the per-side markup from the Poisson arrival intensity, which is what makes the total spread decompose into an inventory-risk term that vanishes as the horizon closes and an order-arrival term that does not. The optimal spread works out to $\delta \approx \$1.69$, giving

Inventory-aware quotes: bid = $r - \delta/2 \approx \$97.15$, ask = $r + \delta/2 \approx \$98.85$,

Naive quotes (ignoring inventory): bid = $\$99.15$, ask = $\$100.85$,

the inventory-aware quotes both exactly \$2.00 lower, precisely the reservation-price adjustment, making the market maker's ask more attractive to a potential buyer (encouraging trades that reduce the long position) and its bid less attractive to a potential seller (discouraging trades that would increase it further), all while preserving the same \$1.69 spread width and therefore the same expected compensation per completed round-trip trade.

This inventory-skewing behavior connects directly to the economic decomposition of the bid-ask spread introduced in Chapter 3: the inventory-holding cost component of the spread is precisely the compensation a market maker requires for the risk that its current inventory, positive or negative, moves against it before an offsetting trade arrives, and the Avellaneda-Stoikov framework makes that intuition mathematically precise by deriving the specific optimal quote adjustment as a function of current inventory, remaining time horizon, and the asset's volatility. Notice, too, that the skew grows with γ , σ^2 , and the size of the inventory position q itself: a more risk-averse market maker, a more volatile asset, or a larger accumulated position all call for a more aggressive skew, exactly the intuition that inventory risk becomes more urgent to offload precisely when it is larger or more dangerous to hold. A market maker who ignores inventory risk entirely, quoting a fixed spread around fair value regardless of current position, is exposed to exactly the kind of accumulating, unhedged directional risk that a well-designed inventory model exists to control, an application-specific instance of the broader risk management discipline developed in Chapter 8.

1.10 Automated Market Makers: On-Chain Liquidity Without an Order Book

Every market-making model examined so far, quote-driven dealer markets in Chapter 3, the Avellaneda-Stoikov inventory model just above, presumes a human or algorithmic market maker who actively chooses prices. Decentralized cryptocurrency exchanges built on blockchains such as Ethereum popularized a fundamentally different mechanism: the *automated market maker* (AMM), a smart contract that quotes prices algorithmically from a fixed formula rather than from a human's order placement decisions, and that draws its liquidity from any user willing to deposit assets into a shared pool rather than from a single designated dealer. Hayden Adams launched the first widely used implementation, Uniswap, in November 2018, building on an idea for a formula-driven, order-book-free exchange sketched earlier by Alan Lu at Gnosis and Vitalik Buterin; Uniswap and its many successors (Curve, Balancer, SushiSwap, and the AMMs embedded in most decentralized exchanges) now route billions of dollars of daily trading volume without any party analogous to the specialists or market makers of Chapter 3 ever choosing a price by hand.

The simplest and most widely deployed AMM design is the *constant-product market maker*. A liquidity pool holds reserves of two assets, x units of asset X and y units of asset Y , and the contract enforces that every trade must leave the product of the two reserves unchanged,

$$x \cdot y = k,$$

for a constant k fixed at whatever level the pool's reserves implied when the trade began. The pool's marginal, quoted price of X in terms of Y is simply the ratio of reserves, $p = y/x$: a trader who deposits Δx units of X into the pool must, by the constant-product rule, withdraw

$$\Delta y = y - \frac{k}{x + \Delta x} = y - \frac{xy}{x + \Delta x}$$

units of Y , and because Δy is a nonlinear, concave function of Δx , every trade moves the pool's own quoted price against the trader placing it, exactly the price-impact phenomenon of the Kyle model in Section 1.6, but generated mechanically by a fixed formula instead of an information-asymmetry argument.

A worked example makes the mechanics concrete. Consider a pool holding $x_0 = 100$ ETH and $y_0 = 200,000$ USDC, so $k = 100 \times 200,000 = 20,000,000$ and the pool's marginal price is $p_0 = y_0/x_0 = \$2,000$ per ETH. A trader sells $\Delta x = 10$ ETH into the pool. The new ETH reserve is $x_1 = 110$, so the new USDC reserve must satisfy $x_1 y_1 = k$, giving $y_1 = 20,000,000/110 \approx \$181,818.18$, and the trader receives

$$\Delta y = 200,000 - 181,818.18 = \$18,181.82.$$

The trader's *effective* price on this trade is $18,181.82/10 \approx \$1,818.18$ per ETH, already below the \$2,000 starting price, and the pool's new marginal price after the trade, $y_1/x_1 = 181,818.18/110 \approx \$1,652.89$, has moved even further, since the marginal price reflects the cost of the next, infinitesimally small trade rather than the average cost of the trade just completed.

This gap between the pre-trade marginal price and the effective price paid, here roughly 9.1% of the starting price for a trade equal to 10% of the pool's ETH reserve, is the AMM's version of market impact: unlike the limit order book of Chapter 3, where impact comes from consuming discrete price levels of resting liquidity, an AMM's impact comes from the continuous curvature of $x \cdot y = k$ itself, and it grows sharply, not linearly, as a trade size grows relative to the pool's reserves. This is why real-world swaps of any size specify a maximum acceptable slippage, and why liquidity providers, not traders, are the ones directly exposed to the next risk this section examines.

Anyone can become a liquidity provider (LP) by depositing X and Y in the pool's current ratio and receiving a proportional claim on the pool's future reserves and the small trading fee charged on each swap (typically 0.30% under Uniswap v2's fee schedule).

This is not a passive, riskless activity. Suppose the same LP had deposited assets when the pool was at $x_0 = 100$ ETH, $y_0 = 200,000$ USDC, and the market price of ETH then doubles to \$4,000. Arbitrageurs will trade against the pool until its own marginal price matches the external market, which for a constant-product pool means the new reserves satisfy both $x_1 y_1 = k$ and $y_1/x_1 = p_1 = \$4,000$, giving

$$x_1 = \sqrt{k/p_1} = \sqrt{20,000,000/4,000} \approx 70.7107 \text{ ETH}, \quad y_1 = \sqrt{k p_1} = \sqrt{20,000,000 \times 4,000} \approx \$282,842.71.$$

The LP's share of the pool is now worth $x_1 p_1 + y_1 = 70.7107(4,000) + 282,842.71 \approx \$565,685.42$. Compare this to simply holding the original 100 ETH and \$200,000 USDC without ever depositing into the pool, worth $100(4,000) + 200,000 = \$600,000$ at the same new price.

The LP's position is worth \$34,314.58 less than not providing liquidity at all, a shortfall of 5.72%, a cost universally known in the DeFi literature as *impermanent loss* even though, once the price has actually moved, the loss is entirely realized the moment the LP withdraws ("impermanent" only in the sense that it partially reverses if the price returns to where it started). For a price ratio $r = p_1/p_0$ between the exit and entry price, this loss relative to simply holding the two assets has the closed form

$$\text{IL}(r) = \frac{2\sqrt{r}}{1+r} - 1,$$

which for $r = 2.0$ gives $\text{IL} = 2\sqrt{2}/3 - 1 \approx -5.72\%$, matching the direct calculation above exactly. The formula is symmetric in a specific sense, $\text{IL}(r) = \text{IL}(1/r)$, so a price that halves costs the LP exactly as

much, in percentage terms, as a price that doubles; and it is always non-positive, so a constant-product LP never outperforms simply holding the two assets on price movement alone.

The LP's compensation for accepting this is the trading fee revenue collected on every swap that passes through the pool, and whether that fee income exceeds the impermanent loss over the LP's holding period is the central profitability question of on-chain liquidity provision, structurally analogous to whether Section 1.9's spread income compensates a dealer for the inventory risk it bears. The economic parallel to Section 1.9 runs deeper than a surface analogy. An AMM liquidity pool is a market maker with no discretion: it cannot skew its quotes based on inventory the way the Avellaneda-Stoikov reservation price does, cannot refuse a trade it judges informationally toxic the way a dealer facing suspected insider order flow might, and cannot widen its spread ahead of an anticipated news event. Its only lever is the constant-product curve's built-in price-impact schedule and the flat fee charged per trade; every other risk-management tool developed for human and algorithmic market makers in this chapter and Chapter 3, adverse-selection screening, inventory skewing, quote withdrawal, is unavailable to it by design. This rigidity is precisely what makes the mechanism trustless and permissionless, no counterparty risk, no credit check, anyone can supply or withdraw liquidity at any time. It is also why AMM liquidity pools are structurally more exposed to informed order flow than a discretionary dealer: an arbitrageur who observes that the external market price has moved away from the pool's quoted price can trade against the stale quote with certainty, extracting value that a human market maker, watching the same news, would have already repriced away. This dynamic, formalized in the DeFi literature as *loss-versus-rebalancing*, is a direct, mechanized analog of the adverse-selection component of the bid-ask spread from Chapter 3, and it is the deeper economic source of the impermanent loss quantified above: the LP is, in effect, perpetually offering a slightly stale quote that sophisticated arbitrageurs are always the first to trade against.




Figure 1.3: The constant-product curve $x \cdot y = k$. Selling ETH into the pool moves the pool's reserves down and to the right along the curve; the curvature is what generates price impact directly from the trade size, with no order book involved.

Figure 1.3 plots this constant-product curve and the same 10-ETH swap worked through above: the trade moves the pool from (100,200,000) to (110,181,818) along the hyperbola, and the curve's flattening slope as x grows is exactly why proportionally larger trades in ETH terms extract proportionally less USDC per unit sold, the AMM's mechanical analog of the diminishing marginal liquidity a limit order book exhibits as a large order walks deeper into resting quotes.

1.10.1 Tokenized Assets, Stablecoins, and New Settlement Rails

The AMM machinery above assumed the assets being traded already live on a blockchain. The more consequential institutional development is the reverse direction: moving *traditional* assets onto those rails. Asset tokenization means issuing a conventional security, a money-market fund share, a Treasury bill, a bond, as a digital token whose ownership record is the blockchain itself rather than an entry at a custodian or transfer agent, and by the mid-2020s this had moved from concept to production: major asset managers now run tokenized money-market funds whose shares settle on public blockchains, and several government-bond issues have been placed natively in tokenized form.

The economics driving this are the settlement economics of Chapter 3, not speculative ones. A conventional securities trade settles through the multi-day, multi-intermediary chain of clearinghouse, custodian, and transfer agent described there, with each intermediary charging for its role and the delay itself creating the counterparty risk that central clearing exists to manage. A tokenized security can settle in minutes, around the clock, with delivery and payment executed atomically in a single transaction, no failed leg possible, which collapses both the cost and the counterparty-risk window. The catch is that atomic settlement needs both legs on-chain, which is where stablecoins enter: tokens engineered to hold a fixed value (almost always \$1) by holding reserves of cash and short-term Treasuries against every token issued, functioning as the cash leg of on-chain markets. A stablecoin is economically a money-market fund share that has renounced yield in exchange for functioning as money, and the quality of its reserves is exactly the credit question Chapter 9 would ask of any issuer, a question made vivid whenever a stablecoin has “broken the buck” and traded below its peg in a stress episode. Central banks have studied issuing the sovereign version, a central bank digital currency (CBDC), though as of this writing large-economy deployments remain pilots rather than production systems.

For the quantitative reader, the relevance is threefold. Tokenized markets generate exactly the granular, timestamped, public transaction data this chapter’s methods consume, but around the clock and with every trade visible. The market-making and execution problems on these venues are the same Kyle-model and inventory-management problems developed above, transplanted onto AMM mechanics. And the migration itself is an institutional-plumbing shift of the kind Chapter 3 argued quantitative practitioners ignore at their peril, since a settlement cycle collapsing from days to minutes changes collateral needs, intraday liquidity management, and the very meaning of counterparty exposure.

1.11 A Worked Example: A Pairs-Trading Backtest

Chapter 7 introduced pairs trading conceptually using six days of prices for two hypothetical bank stocks, finding a spread mean of $\bar{s} = 2.5$, a standard deviation of $\hat{\sigma}_s = 0.548$, and a final z -score of 0.91, too small to trigger a trade under the common rule of entering when $|z| > 2$. Extending that same series four more days, with new prices for Stock A of \$58, \$56.50, \$54.80, and \$54.50, and for Stock B of \$54, \$53, \$52, and \$52, produces the z -scores in Table 1.4, computed using the mean and standard deviation estimated from the original six days (an out-of-sample application of parameters estimated in-sample, exactly the walk-forward discipline of Chapter 7).

Table 1.4: Out-of-sample spread and z -score for the extended pairs-trading example

Day	7	8	9	10
Spread ($A - B$)	4.00	3.50	2.80	2.50
z -score	2.74	1.83	0.55	0.00

Day 7’s z -score of 2.74 exceeds the entry threshold of 2, triggering a trade: short 1 share of Stock A and long 1 share of Stock B (the $\beta \approx 1$ simple-spread convention Chapter 7 used), betting that the unusually wide spread will revert toward its historical mean. By day 10, the spread has fully reverted to 2.50, almost exactly the in-sample mean, triggering an exit.

The trade’s profit is simply the decline in the spread over the holding period,

$$\text{P\&L} = \text{Spread}_{\text{entry}} - \text{Spread}_{\text{exit}} = 4.00 - 2.50 = \$1.50 \text{ per share-pair,}$$

since shorting the spread profits exactly in proportion to how much it narrows. Relative to the capital required to establish the position (the entry prices of the two legs, $58 + 54 = \$112$), this represents a return of approximately 1.34% over the three-day holding period. This single trade illustrates the complete pairs-trading lifecycle introduced only conceptually in Chapter 7: estimating the relationship in-sample, monitoring it out-of-sample, entering when the spread deviates unusually far from its estimated mean, and exiting once it reverts. One caveat, raised in Chapter 7 and worth repeating here, is that this entire strategy depends on the cointegrating relationship remaining stable. A single successful trade in an illustrative four-day extension is not evidence that the relationship will continue to hold indefinitely.

1.12 Statistical Arbitrage Beyond Pairs

The two-stock pairs trade of Section 1.11 generalizes naturally once more than two assets are involved, using exactly the factor-model machinery of Chapter 12 rather than a simple two-asset regression. Instead of finding a single hedge ratio β between two specific stocks, a statistical arbitrage strategy typically regresses each stock in a universe on a common set of factors, the market, size, and value factors of Chapter 12's multi-factor model, or a larger set of statistically extracted factors from a technique like principal component analysis (introduced in Chapter 6's discussion of dimensionality reduction), and treats the regression residual, the return unexplained by any common factor, as the object to trade. A stock whose residual has drifted unusually far from its own historical mean is a candidate to trade against a basket of other stocks in the same peer group, in exactly the same mean-reversion spirit as the two-stock pairs trade, but diversified across many such residual bets simultaneously rather than concentrated in a single pair.

A concrete example makes this precise. Consider three stocks, W, X, and Y, in the same industry sector, each regressed on the same six months of market excess returns used for the CAPM beta regression in Chapter 12. Using the first five months to estimate each stock's beta and residual volatility, exactly the CAPM regression approach Chapter 12 introduced, Table 1.5 shows each stock's estimated beta and its month-6, out-of-sample residual z -score.

Table 1.5: Estimated betas and out-of-sample residual z -scores for three sector peers

Stock	Estimated β	Month-6 residual	z -score
W	0.98	+1.26%	+13.3
X	1.21	-0.15%	-1.8
Y	0.81	+0.09%	+1.1

Stocks X and Y behave as their historical betas would predict, with residuals well within the range of ordinary in-sample noise. Stock W, by contrast, returned far more than its beta of 0.98 and the month's 3% market move would explain, a residual z -score of over 13, an extraordinarily large deviation that strongly suggests either genuine, stock-specific news (in which case the move may be permanent and should not be traded against) or a temporary, unexplained dislocation (in which case a statistical arbitrage desk would short Stock W against a long position in a market-neutral basket of its sector peers, betting the residual reverts). This is precisely the multi-asset generalization of the two-stock pairs trade in Section 1.11: rather than trading one stock's price against one other specific stock's price, the position is effectively long the sector and short this one stock's excess idiosyncratic move relative to it, diversifiable across many such residual signals computed the same way across an entire universe of stocks simultaneously. This generalization changes the risk profile of the strategy in an important way. A single pairs trade is exposed to the idiosyncratic risk that this particular cointegrating relationship breaks down, exactly the risk flagged in Chapter 7 and reiterated above; a statistical arbitrage book trading hundreds of such residual relationships simultaneously diversifies away much of that single-relationship risk, in the same portfolio-level sense Chapter 12 developed for ordinary asset returns, provided the individual residual bets are not all driven by some common, unmodeled factor that would make them break down together. This last caveat is not hypothetical. Exactly this kind of correlated breakdown across many simultaneously held statistical arbitrage positions was a central feature of the August 2007 "quant quake", when several large, otherwise unrelated quantitative equity funds

experienced simultaneous, severe losses over the course of a few days. It was widely attributed to many funds holding similar statistical arbitrage positions and being forced to unwind them at the same time, causing exactly the kind of crowding-driven, mutually reinforcing price move that Chapter 12's discussion of factor crowding described in the context of long-term portfolio allocation rather than short-term trading.

1.13 Strategy Backtesting: Performance Metrics

Evaluating a backtested strategy, whether the moving-average crossover of Section 1.3 or the pairs trade of Section 1.11, requires more than a single cumulative return figure. The Sharpe ratio, introduced in Chapter 12, applies directly to a strategy's own return series, dividing its mean return by its return volatility; for the moving-average strategy's twelve-day return series (mostly zeros, since the strategy is flat for eight of the eleven return periods), the Sharpe ratio is approximately -0.22 , reflecting the strategy's negative average return over this sample. Maximum drawdown, introduced in Chapter 8, measures the largest peak-to-trough decline in cumulative value over the backtest period, a risk measure that a return-only or even a Sharpe-ratio summary can miss entirely. The moving-average strategy's cumulative return path peaks at a 3.48% gain (after day 7) before falling to a cumulative loss of 8.41% (by day 10), so its maximum drawdown is 8.41%. That is a materially different, and arguably more decision-relevant, number than the final -5.22% cumulative return alone, since it reveals how much value an investor holding the strategy would have seen erode from the strategy's own peak, not just where the strategy ultimately ended up. The hit rate, the fraction of active (non-flat) trading days that were profitable, is 25% for the moving-average strategy, one profitable day out of four active days. It is a useful diagnostic precisely because a strategy can have a positive overall Sharpe ratio with a hit rate below 50% if its winning trades are, on average, sufficiently larger than its losing ones, or vice versa. Hit rate alone, like accuracy alone in Chapter 6's classification metrics, can be a misleading summary of a strategy's economic value when considered in isolation from the size of the gains and losses it produces.

1.13.1 Common Backtesting Pitfalls

Beyond choosing the right performance metrics, a credible backtest must also avoid several data-related traps that can make a strategy look far better than it would have actually performed in real time. Survivorship bias, introduced in Chapter 1 as a general data-quality concern, is especially dangerous in strategy backtesting: testing a strategy only on stocks that are still trading today silently excludes every company that went bankrupt or was delisted, which is precisely the population a well-designed strategy most needs to be tested against, since avoiding future disasters is at least as economically valuable as picking future winners. Look-ahead bias occurs whenever a backtest uses information that would not actually have been available at the time a trading decision was made, for example using a company's finalized, restated financial statements rather than the preliminary figures that were actually public at the time, or, in this chapter's own moving-average example, accidentally using today's signal against today's own return rather than lagging it by one period as Section 1.3 was careful to do. Both biases share the same root cause as the data leakage problem Chapter 6 warned about in a machine learning context: information from outside the genuinely available information set silently improves backtested performance in a way no real-time trading process could ever have replicated, and both require the same discipline to avoid, careful, skeptical attention to exactly what data would have been known, and in what form, at each historical point in time the backtest simulates a decision.

1.14 Implementation Shortfall and Transaction Cost Analysis

The gap between a strategy's theoretical performance, computed from the prices a signal was based on, and its actual realized performance, computed from the prices at which trades were genuinely executed, is called implementation shortfall,

$$\text{Implementation Shortfall} = \text{Execution Price} - \text{Decision Price}.$$

For the TWAP execution in Section 1.7, if a portfolio manager's decision to buy was made when the price stood at \$50.00 (the decision price) but the resulting order was ultimately executed via TWAP at \$50.07, the implementation shortfall is

$$\$50.07 - \$50.00 = \$0.07 \text{ per share} = 14 \text{ basis points,}$$

or \$70 in total on the 1,000-share order. This single number bundles together every friction discussed in this chapter: the market impact of Section 1.5, the timing risk of Section 1.8, and any pure price drift that occurred between the decision and its execution for reasons unrelated to the trade itself. It is the standard metric used in transaction cost analysis (TCA) to evaluate whether a trading desk's execution quality is actually adding or subtracting value relative to the investment decisions it is implementing. A strategy that looks profitable using decision prices alone, exactly the trap Chapter 3 warned against when discussing why market structure matters for quantitative finance, can turn out to be unprofitable once realistic implementation shortfall is subtracted from its theoretical returns. That is why serious backtesting, extending the performance metrics of Section 1.13, should model transaction costs explicitly rather than assume execution at the exact prices a signal was computed from.

1.15 Machine Learning for Trading Signals

Every signal in this chapter so far, the moving-average crossover and the pairs-trading z -score, is a simple, fully interpretable rule with only a handful of parameters. The classification and regression methods of Chapter 6 can, in principle, generate considerably richer signals: a gradient-boosted tree or neural network trained on a large set of engineered features (the lagged returns, rolling volatilities, and volume-based features discussed in Chapter 6's feature engineering section) can potentially capture nonlinear patterns that a fixed-form rule like a moving-average crossover cannot represent at all.

A Head-to-Head Test: Gradient Boosting versus Logistic Regression on a Regime Interaction

The twelve-day price series used throughout the feature-engineering catalog below is far too short to fit any model; testing the claim above properly requires simulating a much longer history. The companion notebook builds 5,000 days of returns with volatility clustering (the same GARCH-style recursion Chapter 7 introduced) and a genuine regime interaction layered on top: 3-day momentum continues, in the trend-following spirit of Section 1.3's crossover strategy, during calm, low-volatility stretches, but reverses into mean-reversion during turbulent, high-volatility stretches, exactly the kind of interaction between two features (momentum and the prevailing volatility regime) that a logistic regression's additive coefficients cannot represent, no matter how it is fit, since a single coefficient on momentum cannot simultaneously be positive in one regime and negative in another. A logistic regression and a gradient-boosted classifier are each trained on the same three engineered features from Table 1.6's catalog, momentum, realized volatility, and RSI, on the first 3,500 simulated days, and evaluated on the final 1,496 days, a genuinely held-out segment the training process never touches. This is consistent with the walk-forward discipline Chapter 7 developed specifically because time-ordered data cannot be split randomly without leaking future information into training.

Logistic regression reaches a held-out AUC of 0.5130, barely above the 0.50 a coin flip would achieve, precisely because its additive structure cannot capture a momentum effect that flips sign depending on the volatility regime. Gradient boosting reaches 0.5328, a modest but genuine 0.0198 improvement that holds up consistently across repeated simulations with different random seeds, confirmed by feature importances of 0.423 (momentum), 0.388 (volatility), and 0.189 (RSI), showing the model is drawing on both momentum and volatility jointly rather than either feature alone. Both AUCs are modest in absolute terms, and deliberately so: financial return data have a genuinely low signal-to-noise ratio, exactly the caution Chapter 6 raised repeatedly, so an AUC in the low 0.50s here is a realistic stand-in for the kind of small, hard-won edge a real trading signal typically offers, not a shortcoming of the exercise. The result is nonetheless a meaningful, honest demonstration of the claim at the top of this section: gradient boosting captures a genuine nonlinear interaction a linear-in-its-features model structurally cannot, even though the resulting edge, like most genuine trading edges, is small rather than dramatic.

A Feature-Engineering Catalog for Trading Signals

Chapter 6’s feature engineering section developed the general financial-ML patterns, ratios, lagged and rolling windows, interaction terms, categorical encodings, that apply across credit scoring, fraud detection, and every other application in this book. Trading signals draw on a narrower, more specialized catalog built almost entirely from price, volume, and order-book history, organized here into four families. **Price-based** features summarize recent direction and magnitude: momentum, the trailing return over a fixed lookback window, and realized volatility, the rolling standard deviation of returns over the same kind of window, are the two most common building blocks, with the moving averages of Section 1.3 a third. **Technical indicators** repackage price history into a bounded, more directly interpretable scale; the Relative Strength Index (RSI), worked through numerically below, is one of the most widely used, alongside the Absolute Price Oscillator (APO), the difference between a fast and a slow exponential moving average, and Bollinger Bands, which wrap a moving average in a pair of volatility-scaled upper and lower bands; both are also worked through numerically below. **Volume-based** features, raw volume, dollar volume (price times volume), and volume relative to its own recent average, capture how much conviction accompanies a price move, a dimension none of the price-based features above can see on their own. **Microstructure-based** features draw directly on Chapter 3 and this chapter’s own market-quality material: the bid-ask spread, and the order-flow imbalance that motivated Kyle’s model in Section 1.6, both change on a timescale of individual trades rather than daily closes and can supply information a purely price-based feature never sees at all.

Table 1.6 computes three of these features, reusing the identical twelve-day price series from Section 1.3’s moving-average crossover example: a 3-day momentum ($Mom3_t = P_t/P_{t-3} - 1$), a 3-day realized volatility (the rolling standard deviation of the three most recent daily returns), and a 3-period RSI, $RSI = 100 - 100/(1 + RS)$ where RS is the average of the window’s positive price changes (gains) divided by the average magnitude of its negative changes (losses).

Table 1.6: Three engineered trading-signal features, computed from Section 1.3’s twelve-day price series

Day	Price	Mom3 (%)	Vol3 (%)	RSI3
4	106	6.00	0.781	100.00
5	110	8.91	0.732	100.00
6	115	11.65	0.667	100.00
7	119	12.26	0.450	100.00
8	118	7.27	2.328	90.00
9	114	−0.87	2.835	44.44
10	109	−8.40	1.493	0.00
11	104	−11.86	0.523	0.00
12	100	−12.28	0.313	0.00

The three columns tell a genuinely complementary story about the same twelve days Section 1.3 already walked through. Momentum simply tracks the trend, rising through day 7 and falling steadily thereafter, exactly mirroring the price series it is computed from. Realized volatility does something a momentum feature cannot: it stays low and calm throughout the smooth days-4-through-7 uptrend (0.45% to 0.78%) and then spikes sharply, to 2.3% and 2.8%, exactly at the days 8 and 9 trend reversal Section 1.3 identified as the crossover strategy’s whipsaw, a concrete instance of the volatility clustering Chapter 7’s GARCH material developed at length: volatility rises around genuine regime changes, not uniformly through time, which is precisely why realized volatility is itself a useful predictive feature and not merely a risk-management input. RSI illustrates a real limitation alongside its usefulness: because every change in the days-4-through-7 window is positive, RSI saturates at exactly 100 for four consecutive days, the indicator’s conventional “overbought” signal, but it cannot distinguish a mild uptrend from an explosive one once every recent change has the same sign, and it falls just as quickly to a saturated 0 (“oversold”) once the reversal takes hold and every recent change turns negative. A model trained only on RSI would have no way to tell days 4 through 7 apart from each other, despite their visibly different momentum, which is exactly why practitioners typically feed a model both a bounded indicator like RSI and the underlying continuous features, momentum and

realized volatility among them, side by side, letting the model itself learn how much weight each deserves rather than discarding the information a saturated indicator throws away.

Two further technical indicators are common enough in practice to be worth working through numerically alongside RSI, both computed from an exponentially weighted moving average (EMA) rather than the simple moving average of Section 1.3: an EMA weights recent prices more heavily than older ones, updating as $EMA_t = \lambda P_t + (1 - \lambda)EMA_{t-1}$ with smoothing factor $\lambda = 2/(N + 1)$ for an N -period window, seeded at the first available price.

The **Absolute Price Oscillator (APO)** is simply the difference between a fast and a slow EMA, $APO_t = EMA_t^{\text{fast}} - EMA_t^{\text{slow}}$; industry practice typically uses a 10-day fast and 40-day slow window (or, in the closely related MACD indicator, a 12-day and 26-day pair smoothed further by a signal line), but Table 1.7 instead uses a 3-day and 6-day pair, scaled down to fit this chapter's twelve-day illustrative series, the same way Section 1.3 scaled its own moving-average windows down from industry-standard lengths.

Table 1.7: Absolute Price Oscillator (3-day and 6-day EMA), computed from the same twelve-day price series

Day	Price	EMA3	EMA6	APO
1	100	100.000	100.000	0.000
2	101	100.500	100.286	0.214
3	103	101.750	101.061	0.689
4	106	103.875	102.472	1.403
5	110	106.938	104.623	2.314
6	115	110.969	107.588	3.381
7	119	114.984	110.849	4.136
8	118	116.492	112.892	3.600
9	114	115.246	113.208	2.038
10	109	112.123	112.006	0.117
11	104	108.062	109.719	-1.657
12	100	104.031	106.942	-2.911

The APO column tells the same trend story as momentum but with an EMA's smoother, weighted-average lens: it is positive and climbing throughout the days-1-through-7 uptrend, peaking at 4.136 exactly at day 7's local high, then falls steadily and turns negative by day 11 once the fast EMA has crossed back below the slow one, a lagging but directly interpretable confirmation of the same trend reversal Section 1.3's crossover strategy was built to detect.

Bollinger Bands take a different approach to the same underlying price series, wrapping a simple moving average in bands set a fixed number of standard deviations above and below it: $Middle_t = SMA_t$, $Upper_t = SMA_t + k \cdot \sigma_t$, and $Lower_t = SMA_t - k \cdot \sigma_t$, where σ_t is the rolling standard deviation of price (not return) over the same window as the moving average and k is conventionally set to 2. Industry practice typically uses a 20-day window; Table 1.8 instead uses a 5-day window for the same reason Table 1.7 shortened its EMA windows.

The band width (upper minus lower) is itself a volatility indicator, and it tells exactly the same story as Table 1.6's Vol3 column told numerically: the bands are tight while days 1 through 5 climb smoothly, then widen sharply from day 6 onward and stay wide through day 12, tracking the same days-8-and-9 volatility spike Vol3 flagged and the same GARCH-style volatility clustering Chapter 7 introduced, only expressed geometrically around the price chart rather than as a single rolling standard deviation. Because both indicators are, at bottom, further transformations of the same rolling mean and rolling dispersion of price already summarized by momentum and realized volatility, neither APO nor Bollinger Bands supplies information a model could not in principle extract from those two features directly; their real value is the packaging, a chart-ready, bounded-relative-to-price representation that a human trader can read at a glance, which is exactly why both remain popular in practitioner tools even though a machine-learning pipeline is free to work from the underlying momentum and volatility features instead.

Every indicator in this catalog so far, momentum, realized volatility, RSI, APO, Bollinger Bands, is computed from a single asset's own price history and looks backward at what has already happened. The

Table 1.8: Bollinger Bands (5-day SMA, $k = 2$), computed from the same twelve-day price series

Day	Price	Middle (SMA5)	Upper	Lower
1	100	100.000	100.000	100.000
2	101	100.500	101.500	99.500
3	103	101.333	103.828	98.839
4	106	102.500	107.083	97.917
5	110	104.000	111.266	96.734
6	115	107.000	117.040	96.960
7	119	110.600	122.234	98.966
8	118	113.600	123.447	103.753
9	114	115.200	121.575	108.825
10	109	115.000	122.043	107.957
11	104	112.800	124.071	101.529
12	100	109.000	122.023	95.977

CBOE Volatility Index (VIX), introduced in Chapter 5 as the market-implied volatility aggregated across S&P 500 index options, supplies a genuinely different kind of feature: a market-wide, forward-looking measure of the entire market's expected volatility over the next 30 days, rather than any single stock's own recent price behavior. Trading strategies routinely use the VIX level, or its rate of change, as a regime filter layered on top of the asset-specific signals developed above: a momentum strategy's whipsaw risk (Section 1.3's own days-8-and-9 reversal is a small-scale illustration) tends to be worse when VIX is elevated and the broad market is already unsettled, so a strategy might reduce position sizes, widen its stop-loss thresholds, or stand aside entirely once VIX crosses a specified threshold, rather than treating every trading day as statistically identical regardless of the prevailing volatility regime. Because VIX is itself directly tradable, through VIX futures, options, and exchange-traded products, some strategies also trade VIX-linked instruments directly as a hedge against the same broad market volatility that erodes the returns of the asset-specific strategies developed throughout this chapter.

Two disciplines from earlier chapters apply to every feature in Table 1.6 with particular force. First, the same point-in-time rule Section 1.3 enforced for the moving-average signal applies to every engineered feature here: a momentum, volatility, or RSI value computed from prices through day t can only be used to trade day $t+1$'s return, never day t 's own, and a backtest that lines up a feature with the same day's return it was partly computed from is committing exactly the data leakage Chapter 6 warned against, only easier to miss because the feature looks like an input rather than an answer. Second, Chapter 7's stationarity material explains why every feature in the table is a *transformation* of price rather than the raw price level itself. A stock's price can trend arbitrarily far from its starting value over the life of a dataset, exactly the non-stationarity Chapter 7's unit-root discussion described, so a model trained on raw price levels from one historical period would need to extrapolate to price ranges it never saw in training. Momentum, RSI, and volume ratios are all normalized to a roughly stationary, comparable scale regardless of the underlying stock's price level or the historical period examined. That is exactly why they, rather than price itself, are the features actually handed to a model.

A quantitative equity manager applying these tools typically does not predict an absolute return target directly but instead trains a model to produce a cross-sectional ranking, a relative ordering of an entire universe of stocks from most to least attractive over the next period, since ranking is often both easier to learn reliably and more directly useful for portfolio construction than a precise point forecast of each stock's return. This ranked output feeds naturally into the mean-variance and risk-parity machinery of Chapter 12: rather than optimizing over raw historical returns, the portfolio optimizer takes the model's ranked or scored predictions as its expected-return input, combining a machine-learning-generated forecast with the risk-management discipline developed for classical portfolio construction, exactly the kind of layered, classical-plus-modern pipeline emphasized throughout this book.

Beyond price and volume history, an increasingly important feature source is order-book-derived microstructure data itself, such as the order-flow imbalance (the net difference between resting buy and sell

volume at the top of the book) that motivated Kyle's model in Section 1.6, and alternative data sources such as those introduced in Chapter 17, satellite imagery, web traffic, or transaction data, which can supply a genuinely independent source of predictive information uncorrelated with the price and volume history every other market participant already observes. The same overfitting and data-leakage concerns Chapter 6 raised apply with particular force in this setting: financial return data are extremely noisy relative to any genuine predictive signal they contain, so a sufficiently flexible model, evaluated carelessly, can produce an excellent-looking backtest that is really fitting historical noise, exactly as the high-degree polynomial in Chapter 6 fit noise in its training data. Walk-forward validation, introduced in Chapter 7 specifically because a random train-test split leaks future information into the training set for time-ordered data, is not optional but essential for any machine-learning-based trading signal, and even a walk-forward-validated signal must still be evaluated net of the implementation shortfall and market-impact costs developed earlier in this chapter before its statistical performance can be translated into a genuine, tradable economic edge.

The methods developed throughout this chapter apply across a wide range of trading horizons, from a multi-day pairs trade to an execution algorithm working an order over a single afternoon. At the shortest end of that spectrum, high-frequency trading pushes every one of this chapter's ideas, order types, market impact, inventory management, down to time horizons of milliseconds or microseconds, where the physical latency of transmitting information and orders becomes a first-order concern rather than a rounding error. High-frequency trading is substantial enough a topic, with its own distinctive technology, market-making techniques, and regulatory debates, that it receives its own dedicated treatment in Chapter 14 rather than being folded into this chapter's more general execution and microstructure framework.

1.16 News and Event-Driven Trading

A distinct signal category, related to but separate from the momentum and mean-reversion signals introduced earlier in this chapter, trades directly on the arrival of new information: an earnings announcement, a central bank rate decision, a merger announcement, or a news headline. Because such events are, by construction, largely unpredictable in their timing (a firm's exact earnings-release date is known in advance, but the content of the announcement is not), event-driven strategies typically focus less on predicting when a signal will arrive and more on reacting to it faster and more accurately than other market participants once it does. This is precisely the domain where the natural language processing methods covered in Chapter 16 intersect directly with the execution machinery of this chapter: extracting a sentiment score or a structured summary from an earnings call transcript or a news wire, as Chapter 16 develops in detail, is only half of an event-driven trading strategy. The other half is exactly this chapter's concern, converting that extracted signal into an actual position quickly enough to matter (since an interpretable but slowly generated signal may arrive well after the price has already moved to reflect the same information) and executing the resulting trade while managing the market impact and implementation shortfall costs developed in Sections 1.5 and 1.8, which are often especially severe immediately after a major news event, when liquidity is thin and many other participants are trying to trade on the same information simultaneously.

1.17 Circuit Breakers and Trading Halts

Chapter 3 described the 2010 Flash Crash, in which the Dow Jones Industrial Average fell and then recovered nearly 1,000 points within minutes, driven in significant part by algorithmic trading systems interacting in ways no individual participant had designed or anticipated. One direct regulatory response to that episode, and to the general risk that automated trading can amplify a price move far beyond what changed fundamentals would justify, is the circuit breaker: an automatic, exchange-wide (or single-stock) trading halt triggered once price moves beyond a specified threshold within a specified time window, giving market participants, human and algorithmic alike, a pause to assess whether a price move reflects genuine new information or a mechanical, self-reinforcing feedback loop of the kind that characterized the Flash Crash.

Circuit breakers interact directly with the execution and market-making concepts developed earlier in this chapter. A market maker managing inventory under the Avellaneda-Stoikov framework of Section 1.9, for example, must decide how to handle a position when trading in a security is suddenly halted, since the model's continuous quote-adjustment logic assumes the ability to trade continuously, an assumption a halt

violates by design. Similarly, an execution algorithm using a TWAP or VWAP schedule spread across a trading day, as introduced earlier in this chapter, must have explicit logic for what to do if trading is halted partway through the schedule, either pausing and resuming once trading reopens or reassessing the remaining execution plan entirely, since simply proceeding as though nothing happened once trading resumes could execute the remaining slices at prices that no longer reflect the market conditions the original schedule was designed around.

1.18 Case Vignette: The Knight Capital Trading Glitch

On August 1, 2012, Knight Capital Group, then one of the largest market makers in US equities, deployed new trading software to its production servers. A deployment error left obsolete code active on one of eight servers, code that, combined with the new software, caused the firm's systems to send a rapid, unintended stream of erroneous orders into the market. Within approximately 45 minutes, before the problem was identified and the system halted, Knight Capital had accumulated large, unintended positions across dozens of stocks, and unwinding them resulted in a loss of about \$440 million, an amount exceeding the firm's entire market capitalization at the time and one that ultimately forced Knight into a rescue merger with a competitor.

This episode is a direct, real-world illustration of the technology and latency risk flagged later in this chapter's challenges section, and it is worth being specific about why it happened, since the failure was not a flawed trading strategy in the sense this chapter has otherwise discussed. Knight's algorithms were not making a bad prediction or a poorly calibrated risk-reward tradeoff; a software deployment error caused the firm's own execution infrastructure, the very machinery this chapter has treated as neutral plumbing for implementing a trading decision, to behave in a way no human ever intended or approved. The speed advantage that makes algorithmic and high-frequency trading profitable, the same speed this chapter's discussion of latency emphasized, is precisely what turned a deployment bug into a \$440 million loss in well under an hour: a human-paced trading operation making the same kind of error would very likely have noticed and intervened long before losses accumulated to anything close to that scale. The lesson generalizes directly to the model risk management framework introduced in Chapter 8: an algorithmic trading system requires the same independent testing, staged deployment, and real-time monitoring and kill-switch capability that a bank's VaR or credit model requires, precisely because the speed that makes automation valuable is the same speed that makes an undetected error catastrophic.

1.19 Challenges in Algorithmic Trading and Market Microstructure

Several challenges recur across nearly every algorithmic trading application.

Backtest overfitting. With enough historical data and enough candidate strategies, some rule will appear profitable purely by chance, the same multiple-testing concern raised in Chapters 6, 7, and 12; a strategy discovered by searching over many parameter combinations (different moving-average window lengths, different z -score thresholds) needs to be validated out of sample considerably more skeptically than one derived from a specific, pre-specified economic hypothesis.

Market impact is itself uncertain and strategy-dependent. The clean market-impact numbers in Section 1.5 come from a simplified, static order book; real market impact depends on prevailing volatility, on whether other market participants can detect and react to a large order being worked, and on overall market liquidity conditions that themselves vary over time, all of which make market impact considerably harder to estimate precisely than the other inputs to a trading strategy.

Latency and technology risk. Execution algorithms and market-making strategies alike depend on fast, reliable technology, and a software bug, a connectivity failure, or a delay in receiving market data can turn a theoretically sound strategy into a source of real, sometimes rapid, losses; this is the specific operational risk category from Chapter 9 applied to a domain where a failure can compound within seconds rather than days.

Adverse selection from faster participants. A trading algorithm that posts a limit order risks that order being filled specifically by a faster or better-informed participant exactly when doing so is disadvantageous, the same adverse-selection component of the bid-ask spread introduced in Chapter 3; the more of this risk a strategy is exposed to, the wider the effective spread it must earn just to break even.

Regulatory and market-structure change. Algorithmic and high-frequency trading operate within a regulatory framework (Regulation NMS and its analogues, introduced in Chapter 3) that itself evolves, and a strategy whose profitability depends on a specific market-structure feature, a particular exchange's fee structure or order-matching rule, for example, can see that edge disappear or reverse if the underlying rules change.

Crowding in statistical arbitrage. Section 1.12's discussion of the 2007 quant quake illustrated a specific, historically documented version of a general hazard. When many trading firms independently discover and trade similar statistical relationships, their positions become correlated in a way none of them can observe directly from their own book alone. A shock or forced unwind at one large fund can then trigger simultaneous, mutually reinforcing losses across many ostensibly unrelated strategies, precisely the crowding risk Chapter 12 raised in the context of long-horizon factor investing, but which can unfold over hours rather than months in a statistical arbitrage context.

Exchange fee structures and rebates. Many exchanges charge a fee to remove liquidity (executing against a resting order) while paying a rebate to add it (posting a resting limit order that another participant later executes against). This maker-taker pricing model creates its own subtle incentives independent of the underlying trading signal. A strategy's profitability can depend materially on whether it typically posts passive limit orders or crosses the spread with aggressive market orders, and two strategies with identical raw trading signals can have very different net profitability purely because of how each interacts with a given exchange's specific fee schedule, an execution-level detail entirely separate from whether the underlying signal is economically sound. A realistic backtest, following the transaction-cost discipline of Section 1.13, should account for these fees and rebates explicitly rather than assuming every trade is executed at a single, cost-free price, since even a few basis points of fee difference, compounded across thousands of trades, can be the difference between a strategy that is genuinely profitable after costs and one that only appears to be. This is a further, concrete illustration of the implementation-shortfall discipline introduced earlier in this chapter: the fee schedule itself is one more component of the gap between a strategy's theoretical and realized performance, no less real for being easy to overlook in an initial backtest.

Capacity constraints on strategy scale. A strategy backtested at a small notional size can look highly attractive on a risk-adjusted basis, but the market impact costs developed in Section 1.5 grow, roughly speaking, faster than linearly with order size, so the same strategy run at ten or a hundred times the capital can see its net-of-cost returns deteriorate sharply, or even turn negative, well before the strategy exhausts the raw statistical edge its backtest identified; this capacity limit is a genuine economic constraint on how much capital any single trading idea, however sound, can actually absorb profitably.

1.20 Key Takeaways

Converting a trading signal into a real position and a real position into an actual profit-and-loss series requires confronting the same market microstructure realities Chapter 3 introduced: the bid-ask spread, the limited depth available at each price level, and the price impact of consuming that depth too quickly. The Kyle model gives that price impact a formal economic interpretation, as compensation for the risk of trading against better-informed counterparties, and TWAP and VWAP provide simple, standard benchmarks for execution quality, while the Almgren-Chriss framework formalizes the tradeoff between market impact and timing risk that this chapter's simpler comparisons illustrated concretely.

Market making extends the spread-capturing intuition of Chapter 3 into an explicit inventory-management problem, and pairs trading, introduced only conceptually in Chapter 7, becomes a complete, backtestable strategy once entry and exit rules are applied to real, out-of-sample data, a logic that extends naturally to statistical arbitrage across many assets using the factor models of Chapter 12.

Evaluating any such strategy requires more than a single return figure: the Sharpe ratio, maximum drawdown, and hit rate each reveal a different dimension of performance, and every one of them must be computed net of realistic transaction costs and implementation shortfall, not against the frictionless prices

a signal was originally computed from, since the gap between theoretical and realized performance is often the difference between a strategy that looks good on paper and one that is actually worth trading.

Beyond the mechanics of any single strategy, this chapter's case vignette and its discussion of circuit breakers underscore a broader point that recurs throughout this book: the technology and market structure surrounding a trading decision are not neutral plumbing but active sources of risk in their own right, deserving the same governance discipline Chapter 8 developed for financial risk models generally.

Finally, this chapter's challenges section is a reminder that a statistically sound backtest is necessary but never sufficient: fees, capacity limits, crowding, and the everyday risk of a software or operational failure all sit between a promising historical result and a genuinely profitable, sustainable trading operation, and the machine learning signals of Chapter 6 and the natural-language and alternative-data sources of Chapters 16 and 17 must clear exactly these same practical hurdles before a sophisticated model becomes a trade that actually makes money net of the frictions this chapter has taken care to make explicit.

The next chapter narrows the focus to the extreme, millisecond-scale end of this same spectrum, where every idea introduced here, order types, market impact, inventory management, backtesting discipline, still applies but is compressed into a timescale where speed itself becomes a primary source of competitive advantage.

1.21 Suggested Exercises

1. Using Table 1.1, verify by hand that the 3-day and 6-day moving averages at day 8 are 117.33 and 111.83, and explain why the strategy signal remains long (1) at that point.
2. Recompute the moving-average crossover strategy's cumulative return using a 2-day short average and a 4-day long average on the same twelve-day price series, and compare the resulting signal timing to the 3-day/6-day version in the text.
3. A 500-share market order walks through an order book with 200 shares available at \$30.00 and additional shares available at \$30.05. Compute the average execution price and the price impact in basis points relative to the best price.
4. Using Table 1.3, compute the TWAP and VWAP execution prices for a 2,000-share order, and explain why they differ.
5. Explain, in your own words, why an inventory-skewing market maker (Section 1.9) adjusts both its bid and ask when it accumulates a large position, rather than adjusting only the side of the quote nearest the current inventory.
6. Using the extended pairs-trading data in Table 1.4, compute the trade's profit and loss if the position were instead closed at day 9 rather than day 10, and compare it to the day-10 exit computed in the text.
7. A strategy has a Sharpe ratio of 1.2 and a maximum drawdown of 18%, while a second strategy has a Sharpe ratio of 0.8 and a maximum drawdown of 5%. Discuss which strategy an investor with a low tolerance for large losses might prefer, and why Sharpe ratio alone would not reveal this consideration.
8. A portfolio manager's decision price is \$100.00 and the resulting order is executed at an average price of \$100.12 on a 5,000-share order. Compute the implementation shortfall in basis points and in total dollars.
9. Using Table 1.6's twelve-day price series, compute the 3-day momentum, 3-day realized volatility, and 3-period RSI at day 3, explaining why fewer of these features are available that early in the series than at day 4 onward.
10. Explain, referencing Table 1.6's RSI column, why an indicator that saturates at 0 or 100 can lose information a model might otherwise use, and describe why feeding a model both RSI and the underlying momentum feature side by side addresses this limitation.

11. Using Table 1.7, verify by hand that day 8's APO is 3.600, and explain why APO does not turn negative until day 11, four days after day 7's price high, rather than exactly at it.
12. Using Table 1.8, explain why the band width at day 7 (Upper minus Lower) is so much larger than at day 2, referencing this section's discussion of Table 1.6's Vol3 column.
13. Explain why walk-forward validation, introduced in Chapter 7, is essential when evaluating a machine-learning-based trading signal, and describe a specific way that a random train-test split could produce a misleadingly optimistic backtest.
14. Pick two of the challenges listed in this chapter and describe a concrete scenario, not already used in the text, in which that challenge could cause an algorithmic trading strategy to lose money despite a statistically sound backtest.
15. Using Table 1.2, verify by hand that Kyle's lambda is approximately 1.95 cents per thousand shares, and interpret what a much larger estimated lambda (say, 10 cents per thousand shares) would suggest about that market compared to the one in the text.
16. Explain the difference between an iceberg order and trading in a dark pool, and describe one advantage and one disadvantage each has relative to posting a single, fully visible limit order.
17. Describe how a market maker using the inventory-skewing logic of Section 1.9 should adjust its behavior if trading in the underlying stock is suddenly halted by a circuit breaker, and explain what risk this halt introduces that continuous trading does not.
18. Using the Avellaneda-Stoikov formulas in Section 1.9, recompute the reservation price, optimal spread, bid, and ask for a market maker with the same parameters as the worked example except that it is short $q = -8$ units instead of long 5, and explain why the resulting skew now points in the opposite direction.
19. Explain why the Knight Capital case vignette is better understood as a technology and governance failure than as a bad trading strategy, and describe one specific model risk management practice from Chapter 8 that, if followed, might plausibly have prevented or limited the loss.
20. A statistical arbitrage fund holds 200 similar mean-reversion positions across a sector. Explain, referencing the 2007 quant quake, why this diversification might provide less protection than the number of positions alone would suggest.
21. Using Table 1.5, explain why Stock W's large residual z -score alone does not tell a trader whether to actually put on the suggested trade, and describe what additional information (beyond the regression output itself) an analyst would want before shorting Stock W against its sector peers.
22. A backtest of a value strategy is run only on stocks currently listed on a major exchange, using their full, publicly available price history. Identify the specific bias this introduces, referencing this chapter's discussion of backtesting pitfalls, and describe one concrete step a researcher could take to correct it.
23. Using the Almgren-Chriss worked example, compute κ and the resulting optimal trajectory x_j for the same 250-share order and $N = 5$ intervals if risk aversion were $\lambda = 2$ instead of 0.5 or 5, and state whether the resulting schedule is more or less front-loaded than both trajectories computed in the text.
24. Explain, in your own words, why the fully aggressive single-shot execution in the Almgren-Chriss worked example has zero timing-risk variance, and why this alone does not make it the risk-adjusted optimal strategy even for a highly risk-averse trader.
25. **Challenge (real data).** Download at least two years of daily prices for a real, liquid stock (via `yfinance`) and implement this chapter's moving-average crossover strategy (Table 1.1) using a 20-day short window and a 50-day long window. (a) Backtest the strategy's cumulative return against a simple buy-and-hold benchmark over the full sample, and compute both strategies' Sharpe ratios. (b) Using walk-forward validation rather than a single in-sample backtest, as this chapter's own discussion

(building on Chapter 7) recommends, split your sample into several sequential training/testing windows and report whether the strategy's out-of-sample Sharpe ratio holds up as well as its in-sample Sharpe ratio. (c) Compute the feature set in Table 1.6 (momentum, realized volatility, RSI) for your real price series, and discuss whether the same volatility-clustering pattern this chapter's toy twelve-day example showed around its crossover reversal appears at your strategy's actual entry and exit points.

1.22 Further Reading

- Adams, Zinsmeister, and Robinson, “Uniswap v2 Core,” 2020. The whitepaper for the constant-product automated market maker formalized in Section 1.10.
- Almgren and Chriss, “Optimal Execution of Portfolio Transactions,” *Journal of Risk*. The original formalization of the market-impact-versus-timing-risk tradeoff developed in Section 1.8.
- Appel, *Technical Analysis: Power Tools for Active Investors*. Written by the creator of the MACD indicator, the closely related, signal-line-smoothed cousin of the Absolute Price Oscillator used in Section 1.15's feature-engineering worked example.
- Avellaneda and Stoikov, “High-Frequency Trading in a Limit Order Book,” *Quantitative Finance*. The original inventory-based market-making model summarized in Section 1.9.
- Bollinger, *Bollinger on Bollinger Bands*. Written by the creator of the Bollinger Bands indicator used in Section 1.15's feature-engineering worked example.
- Chan, *Algorithmic Trading: Winning Strategies and Their Rationale*. A practitioner-oriented treatment of strategy development and backtesting, including the pairs-trading and mean-reversion strategies extended in this chapter.
- Harris, *Trading and Exchanges: Market Microstructure for Practitioners*. A comprehensive treatment of order types, market structure, and execution that extends this chapter's introductory coverage and connects directly back to Chapter 3.
- Kissell, *The Science of Algorithmic Trading and Portfolio Management*. Covers transaction cost analysis and implementation shortfall, introduced in this chapter, in considerably greater quantitative depth.
- Kyle, “Continuous Auctions and Insider Trading,” *Econometrica*. The original 1985 paper formalizing the price-impact model summarized in Section 1.6.
- Wilder, *New Concepts in Technical Trading Systems*. The original source introducing the Relative Strength Index used in Section 1.15's feature-engineering worked example.
- Guéant, Lehalle, and Fernandez-Tapia, “Dealing with the Inventory Risk: A Solution to the Market Making Problem,” *Mathematics and Financial Economics*. A rigorous, closed-form treatment extending the Avellaneda-Stoikov inventory model of Section 1.9 to more realistic order-arrival assumptions.
- Patterson, *Dark Pools: The Rise of the Machine Traders and the Rigging of the U.S. Stock Market*. A journalistic but well-researched account of the rise of dark pools, algorithmic trading, and the market-structure debates touched on in this chapter's discussion of order routing and venue choice.